

Develop PD model, by Brent Oeyen

The goal of this code is to compare the following Probability of Default models

- An IRLS logistic regression with a classic transformation of variable;
- An elastic net logistic regression with raw numeric input variables, calibrated with scipy minimize;
- A Keras DNN binary classification algorithm with raw numeric input variables transformed with a standard score; and
- A Keras elastic net with raw numeric input variables transformed with a standard score.

The following packages are used in the implementation:

```
In [1]: import os
import pandas as pd
import numpy as np
import scipy.optimize as optimize
import scipy.stats as st
from keras.models import Sequential
from keras.layers import Dense
from keras.wrappers.scikit_learn import KerasClassifier
from keras.layers import Dropout
from keras.constraints import maxnorm
from keras.optimizers import SGD
from sklearn.model_selection import cross_val_score
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import StratifiedKFold
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.calibration import CalibratedClassifierCV
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from Codes.PD.PD_tests import *
from Codes.PD.Logistic_regression import *
```

Using TensorFlow backend.

Load dataset, source: "<http://archive.ics.uci.edu/ml/machine-learning-databases/statlog/german/>
(<http://archive.ics.uci.edu/ml/machine-learning-databases/statlog/german/>)".

```
In [2]: location = os.getcwd() + '/Data/german data numeric'
dataframe = pd.read_csv(location, header=None, delimiter=r"\s+")
```

Split data set into input (X) input variables and binary (Y) output variable:

```
In [3]: X = dataframe.values[:,0:24].astype(float)
dataframe.iloc[:, 24] = dataframe.iloc[:, 24]==2
Y = dataframe.values[:,24].astype(float)
```

Transform input variables with a logit transformation per quantiles:

```
In [6]: def transform(df, j):
df.columns      = ['var', 'binary']
try:
    df['decile'] = pd.qcut(df['var'], q=np.arange(1, 51)/50, labels=False)
    df          = pd.merge(df.astype(float), df.groupby('decile', as_index=False)
['binary'].mean(), \
                           on='decile', how='inner')
except:
    df          = pd.merge(df.astype(float), df.groupby('var'      , as_index=False)
['binary'].mean(), \
                           on='var'      , how='inner')
pi              = df.iloc[:, -1].fillna(0.5).replace(0, 0.00001)
return pd.DataFrame({'X'+str(j): np.log(pi/(1-pi)).replace(np.inf, 15)})
X_scaled       = transform(dataframe.iloc[:, [0,24]], 1)
for i in range(1, dataframe.shape[1]-1):
    X_scaled['X'+str(i+1)] = transform(dataframe.iloc[:, [i,24]], i+1)
X_scaled       = X_scaled.fillna(0)
X_train, X_test, Y_train, Y_test = train_test_split(X_scaled, Y, test_size=0.2
, random_state=42)
X_tr   , X_tes  , Y_tr   , Y_tes  = train_test_split(X      , Y, test_size=0.2
, random_state=42)
```

Output variance of each transformed input variable as an indication of which variables contain a lot of information:

```
In [7]: print(X_train.var())
```

```
X1      0.762004
X2      3.421997
X3      0.295841
X4     14.081804
X5      0.226596
X6      0.085566
X7      0.043210
X8      0.003601
X9      0.110774
X10     1.683886
X11     0.053565
X12     0.013477
X13     0.000043
X14     0.006430
X15     0.056956
X16     0.041302
X17     0.065301
X18     0.000003
X19     0.014833
X20     0.036471
X21     0.079953
X22     0.000149
X23     0.002296
X24     0.000869
dtype: float64
```

Verify whether the transformation of variable was successful by comparing the rank correlation:

```
In [8]: print("Raw input variable X2: ", st.kendalltau(Y, X[:, 1]))
print('Transformed input variable X2:', st.kendalltau(Y, X_scaled.iloc[:, 1]))

Raw input variable X2: KendalltauResult(correlation=0.17609245525504136, pvalue=7.975280722434196e-11)
Transformed input variable X2: KendalltauResult(correlation=0.013532833811191117, pvalue=0.6174826267368242)
```

Create a Master Scale of 22 ratings mapped to equidistant PDs in logit space:

```
In [9]: ratings = 22 #Number of rating grades
PD_min = 0.0003 #Minimum PD value (regulatory threshold)
slope = (np.log((2-PD_min)/PD_min)-0)/(ratings-1) #Slope between min PD value and default in logit space
MS = 2/(1+np.exp(slope*pd.Series(list(range(ratings))))
idx = pd.IntervalIndex.from_arrays(MS[1:].append(pd.Series(0)),
MS, closed='left')
```

IRLS logistic regression with a classic transformation of variable

Calibrate model:

```
In [10]: betas_start = np.append(np.log(Y_train.mean()/(1-Y_train.mean()
)), np.zeros(8))
betas_IRLS, y_train_IRLS = logistic_regression().IRLS(betas_start, np.append(np.ones([len(X_train.index), 2]), \
X_train.iloc[:, np.r_[0:4, 7:9, 16, 20]], axis=1)
[:, 1:], Y_train)
X_IRLS = np.append(np.ones([len(X_scaled.index), 2]), \
X_scaled.iloc[:, np.r_[0:4, 7:9, 16, 20]], axis=1)
[:, 1:]
IRLS_all = pd.DataFrame({'PD': 1/(1+np.exp(X_IRLS.dot(-betas_IRLS)))})
IRLS_all['Y'] = Y
IRLS_all['rating_PD'] = MS[idx.get_indexer(IRLS_all.PD)].values
```

IRLS converged after 4 iterations.

Output results for a regression with the 8 highest correlated variables:

```
In [11]: print("Performance metrics: [AUC all %.2f%%]" % (PD_tests().AUC(IRLS_all.Y, IRLS_all.rating_PD,0)[0]*100))
print("Coefficients regression:", betas_IRLS)
dummy = PD_tests().Jeffrey(IRLS_all, 'rating_PD', 'PD', 'Y')
print('Jeffrey test')
print(dummy.iloc[:, np.r_[0, 3, 6, 10:12]])
```

```
Performance metrics: [AUC all 53.59%]
Coefficients regression: [ 0.63104141 -0.00249486 -0.02409827  0.21552432 -0.00565019  0.70273931
 0.23757129  0.47568141  0.12372487]
```

```
Jeffrey test
```

	rating_PD count	PD mean	Y mean	H0	p_val
rating_PD					
0.21890161217657667	2.0	0.198208	0.500000	0.198208	0.140557
0.31496223179993454	529.0	0.273935	0.270321	0.273935	0.571111
0.4426995169025332	469.0	0.333599	0.332623	0.333599	0.515712
Portfolio	1000.0	0.805742	1.102944	0.301766	0.546622

Given the poor performance of the scaled input variables, the raw input variables are used instead:

```
In [20]: betas_start = np.append(np.log(Y_tr.mean()/(1-Y_tr.mean())) , np.zeros(8))
betas_IRLS, y_train_IRLS = logistic_regression().IRLS(betas_start, np.append(np.ones([len(X_tr), 2]), \
X_tr[:, np.r_[0:4, 7:9, 16, 20]], axis=1)[: , 1:], Y_tr)
X_IRLS = np.append(np.ones([len(X), 2]), \
X[:, np.r_[0:4, 7:9, 16, 20]], axis=1)[: , 1:]
IRLS_all = pd.DataFrame({'PD': 1/(1+np.exp(X_IRLS.dot(-betas_IRLS)))})
IRLS_all['Y'] = Y
IRLS_all['rating_PD'] = MS[idx.get_indexer(IRLS_all.PD)].values
print("Performance metrics: [AUC all %.2f%%]" % (PD_tests().AUC(IRLS_all.Y, IRLS_all.rating_PD,0)[0]*100))
print("Coefficients regression:", betas_IRLS)
dummy = PD_tests().Jeffrey(IRLS_all, 'rating_PD', 'PD', 'Y')
print('Jeffrey test')
print(dummy.iloc[:, np.r_[0, 3, 6, 10:12]])
```

```
IRLS converged after 5 iterations.
Performance metrics: [AUC all 77.88%]
Coefficients regression: [ 0.77847929 -0.56161866  0.0246094 -0.36420994  0.00514031 -0.0544452
 0.20944244 -1.17560474 -0.51219642]
```

```
Jeffrey test
```

	rating_PD count	PD mean	Y mean	H0	p_val
rating_PD					
0.029760560578478736	4.0	0.027608	0.000000	0.027608	0.374643
0.04491339548733918	26.0	0.039108	0.000000	0.039108	0.852174
0.06751703372078455	63.0	0.056960	0.031746	0.056960	0.800781
0.10090929621852611	96.0	0.083958	0.083333	0.083958	0.488103
0.14953859341903025	115.0	0.125885	0.130435	0.125885	0.428017
0.21890161217657667	135.0	0.184365	0.177778	0.184365	0.569288
0.31496223179993454	162.0	0.263878	0.253086	0.263878	0.617424
0.4426995169025332	163.0	0.377768	0.404908	0.377768	0.236454
0.6036825140620321	128.0	0.526771	0.531250	0.526771	0.460227
0.7933816324488658	90.0	0.688752	0.677778	0.688752	0.594176
1.0	18.0	0.850305	0.833333	0.850305	0.609116
Portfolio	1000.0	3.225356	3.123647	0.301186	0.530763

Elastic net logistic regression with raw numeric input variables

```
In [55]: solution = optimize.minimize(fun=logistic_regression().el_logicreg
, x0=betas_IRLS, args=(Y_tr, \
                        X_tr[:, np.r_[0:4, 7:9, 16, 20]], 0.3, 0.5), method='Nelder-Mead', options={"maxiter":5000})
X_LL = np.append(np.ones([len(X), 2]), X[:, np.r_[0:4, 7:9, 16, 20]], axis=1)[: , 1:]
LL_all = pd.DataFrame({'PD': 1/(1+np.exp(X_LL.dot(-solution.x))))})
LL_all['Y'] = Y
LL_all['rating_PD'] = MS[idx.get_indexer(LL_all.PD)].values
```

Output results for a regression with the 8 highest correlated variables:

```
In [57]: print("Results logistic regression ML: Elastic net (Lambda=0.3, L1 ratio=0.5)")
print("Performance metrics: [AUC all %.2f%%]" % (PD_tests().AUC(LL_all.Y, LL_all.rating_PD,0)[0]*100))
print("Coefficients regression:", solution.x)
dummy = PD_tests().Jeffrey(LL_all, 'rating_PD', 'PD', 'Y')
print('Jeffrey test')
print(dummy.iloc[:, np.r_[0, 3, 6, 10:12]])
```

Results logistic regression ML: Elastic net (Lambda=0.3, L1 ratio=0.5)

Performance metrics: [AUC all 72.35%]

Coefficients regression: [1.34782851e+00 -3.40238037e-01 5.28438161e-02 -4.65549331e-01 -1.58887734e-02 -1.50196892e-01 -1.39382220e-04 9.86211126e-10 -1.51418317e-01]

Jeffrey test

	rating_PD count	PD mean	Y mean	H0	p_val
rating_PD					
0.04491339548733918	3.0	0.036647	0.333333	0.036647	0.023039
0.06751703372078455	5.0	0.063221	0.400000	0.063221	0.009732
0.10090929621852611	28.0	0.091162	0.035714	0.091162	0.849184
0.14953859341903025	87.0	0.123939	0.080460	0.123939	0.895987
0.21890161217657667	162.0	0.185623	0.154321	0.185623	0.847680
0.31496223179993454	184.0	0.263588	0.228261	0.263588	0.862264
0.4426995169025332	204.0	0.376796	0.225490	0.376796	0.999998
0.6036825140620321	195.0	0.514633	0.482051	0.514633	0.818697
0.7933816324488658	119.0	0.685217	0.605042	0.685217	0.968352
1.0	13.0	0.841622	0.769231	0.841622	0.776087
Portfolio	1000.0	3.182449	3.313904	0.362034	0.999982

Keras DNN with raw numeric input variables transformed with a standard score

```

In [58]: def create_binary():
            # create model
            model = Sequential()
            model.add(Dense(24, input_dim=24, activation='relu'))
            model.add(Dropout(0.2))
            model.add(Dense(24, activation='relu'))
            model.add(Dropout(0.2))
            model.add(Dense(1, activation='sigmoid'))
            # Compile model
            model.compile(loss='binary_crossentropy', optimizer='adam', metrics=[
'accuracy'])
            return model
estimators = []
estimators.append(('standardize', StandardScaler()))
estimators.append(('mlp', KerasClassifier(build_fn=create_binary, epochs=80, batch_size=16, verbose=0)))
pipeline = Pipeline(estimators)
pipeline.fit(X_tr, Y_tr)
kfold = StratifiedKFold(n_splits=10, shuffle=True)
results = cross_val_score(pipeline, X, Y, cv=kfold)
b_pred = pd.DataFrame(pipeline.predict_proba(X_tes))
b_pred.columns = ['ID', 'PD']
b_pred['Y'] = Y_tes
b_pred['rating_PD'] = MS[idx.get_indexer(b_pred.PD)].values
b_all = pd.DataFrame(pipeline.predict_proba(X))
b_all.columns = ['ID', 'PD']
b_all['Y'] = Y
b_all['rating_PD'] = MS[idx.get_indexer(b_all.PD)].values

```

Output results with all input variables being used:

```
In [59]: print("Performance metrics: Mean Accuracy CV %.2f%% (STD %.2f%%) [AUC test %.2f%%] [AUC all %.2f%%]" \
% (results.mean()*100, results.std()*100, PD_tests().AUC(b_pred.Y, b_pred.rating_PD,0)[0]*100, \
PD_tests().AUC(b_all.Y, b_all.rating_PD,0)[0]*100))
dummy = PD_tests().Jeffrey(b_all, 'rating_PD', 'PD', 'Y')
print('Jeffrey test')
print(dummy.iloc[:, np.r_[0, 3, 6, 10:12]])
```

Performance metrics: Mean Accuracy CV 74.90% (STD 2.88%) [AUC test 78.31%] [AUC all 92.09%]

Jeffrey test

	rating_PD count	PD mean	Y mean	H0	p_val
rating_PD					
0.00029999999999999976	27.0	0.000120	0.000000	0.000120	0.064528
0.0004562208354469178	5.0	0.000394	0.000000	0.000394	0.051319
0.0006937632793232519	14.0	0.000580	0.000000	0.000580	0.102292
0.0010549227058496641	15.0	0.000922	0.000000	0.000922	0.133248
0.0016039437457623122	14.0	0.001259	0.000000	0.001259	0.150303
0.002438347229665251	24.0	0.001964	0.041667	0.001964	0.007378
0.0037060190379897217	20.0	0.003030	0.000000	0.003030	0.274099
0.005630882685189334	19.0	0.004764	0.052632	0.004764	0.019101
0.00855121588022675	28.0	0.007344	0.035714	0.007344	0.061640
0.012976261070719905	23.0	0.010860	0.043478	0.010860	0.080403
0.01966854325104041	30.0	0.016925	0.033333	0.016925	0.202517
0.029760560578478736	37.0	0.025303	0.000000	0.025303	0.832971
0.04491339548733918	44.0	0.036913	0.045455	0.036913	0.338925
0.06751703372078455	67.0	0.055981	0.059701	0.055981	0.416403
0.10090929621852611	63.0	0.084232	0.031746	0.084232	0.947688
0.14953859341903025	60.0	0.124943	0.133333	0.124943	0.403778
0.21890161217657667	79.0	0.180652	0.088608	0.180652	0.988616
0.31496223179993454	79.0	0.266653	0.215190	0.266653	0.850006
0.4426995169025332	84.0	0.373171	0.416667	0.373171	0.204085
0.6036825140620321	82.0	0.516930	0.609756	0.516930	0.045744
0.7933816324488658	87.0	0.695557	0.885057	0.695557	0.000015
1.0	99.0	0.885664	0.939394	0.885664	0.037534
Portfolio	1000.0	3.294161	3.631731	0.277586	0.057593

Keras elastic net with raw numeric input variables transformed with a standard score

```
In [63]: LR = LogisticRegression(C=0.3, penalty='elasticnet', solver=
'saga', l1_ratio=0.5, \
max_iter=1000, tol=0.001)
scaler = StandardScaler()
scaler.fit(X_tr)
LR.fit(scaler.transform(X_tr), Y_tr)
LR_pred = pd.DataFrame(LR.predict_proba(scaler.transform(X_tes)))
LR_pred.columns = ['ID', 'PD']
LR_pred['Y'] = Y_tes
LR_pred['rating_PD'] = MS[idx.get_indexer(LR_pred.PD)].values
LR_all = pd.DataFrame(LR.predict_proba(scaler.transform(X)))
LR_all.columns = ['ID', 'PD']
LR_all['Y'] = Y
LR_all['rating_PD'] = MS[idx.get_indexer(LR_all.PD)].values
results = cross_val_score(LR, X, Y, cv=kfold)
```

Output results with all input variables being used:

```
In [64]: print("Results logistic regression:Elastic net (Lambda=0.3, L1 ratio=0.5)")
print("Performance metrics: Mean Accuracy CV %.2f%% (STD %.2f%%) [AUC test %.2f%%] [AUC all %.2f%%]" \
% (results.mean()*100, results.std()*100, PD_tests().AUC(LR_pred.Y, LR_pred.rating_PD,0)[0]*100, \
PD_tests().AUC(LR_all.Y, LR_all.rating_PD,0)[0]*100))
print("Coefficients regression:", LR.coef_)
dummy = PD_tests().Jeffrey(LR_all, 'rating_PD', 'PD', 'Y')
print('Jeffrey test')
print(dummy.iloc[:, np.r_[0, 3, 6, 10:12]])
```

```
Results logistic regression: Elastic net (Lambda=0.3, L1 ratio=0.5)
Performance metrics: Mean Accuracy CV 75.10% (STD 2.84%) [AUC test 80.51%] [AUC all 81.15%]
Coefficients regression: [[-0.65714083  0.30592054 -0.39515111  0.15154385 -0.23556821 -0.15974222
-0.11155633  0.01905668  0.16390315 -0.25583818 -0.22045609  0.14207899
0.          -0.07077531 -0.21050127  0.28600726 -0.2452701  0.20259282
0.10702668  0.01964218 -0.19007338 -0.10013797 -0.04893568 0.          ]]
Jeffrey test
```

	rating_PD count	PD mean	Y mean	H0	p_val
rating_PD					
0.00855121588022675	1.0	0.006156	0.000000	0.006156	0.099792
0.012976261070719905	2.0	0.012371	0.000000	0.012371	0.187656
0.01966854325104041	6.0	0.016613	0.000000	0.016613	0.352909
0.029760560578478736	21.0	0.024908	0.000000	0.024908	0.699527
0.04491339548733918	40.0	0.036768	0.000000	0.036768	0.917541
0.06751703372078455	61.0	0.056392	0.049180	0.056392	0.564286
0.10090929621852611	91.0	0.084324	0.065934	0.084324	0.725217
0.14953859341903025	105.0	0.122635	0.142857	0.122635	0.256485
0.21890161217657667	145.0	0.182612	0.151724	0.182612	0.831874
0.31496223179993454	130.0	0.261591	0.253846	0.261591	0.573564
0.4426995169025332	128.0	0.373937	0.359375	0.373937	0.630607
0.6036825140620321	134.0	0.518178	0.567164	0.518178	0.128114
0.7933816324488658	106.0	0.687456	0.735849	0.687456	0.140659
1.0	30.0	0.843433	0.700000	0.843433	0.978151
Portfolio	1000.0	3.227374	3.025930	0.302073	0.555003

Conclusion

Based on the results of the credit data set, transforming the input data wasn't successful on a small sample population (1K). The imbalance of the observed defaults and performing credit did not create calibration issues.

The Keras library is straight forward to use and requires little development from the user. Albeit fine tuning the configuration parameters to run the models will require some time and experience with both the tool and modelling a given dependent variable. In addition, the Keras library outperforms a cinch implementation of the logistic regression's IRLS and ML elastic net method. The DNN model outperforms the logistic regression based on the AUC of the total population (training and test data) but doesn't on the training set. In addition, the Cross Validation (CV) score of the training set is comparable while the logistic regression produces more stable results.

Considering the comparable performance between DNN and elastic net and the apprehensible results of the elastic net in contrast of the complexity of understanding the probability calculation of the DNN algorithm (not covered in this example), logistic regression with an elastic net is a preferred option for PD models.